

Advanced Out-of-Distribution Detection Leveraging Vision Transformer with Adversarial Attack

JongHyeok Ahn, Leesang Youn, Minho Song

MS. Candidate

University of Seoul

Contents



- Introduction
- Related works
- Method
- Experiment

Limitation of prior Out-of-Distribution Detection

Introduction

Challenged of Near-OOD detection

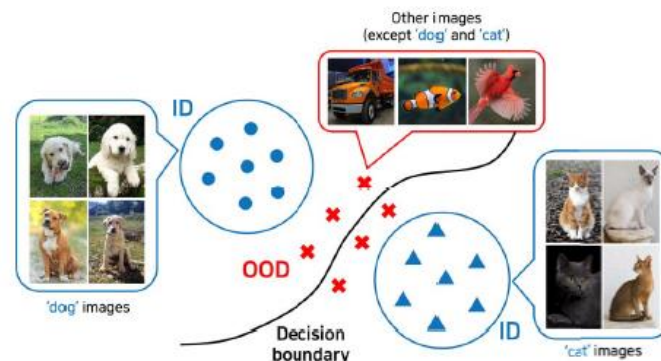
- Near-OOD data is much more similar to ID data to make it harder for existing methods
- Distinguishing between CIFAR-10 and CIFAR-100 is challenging Near-OOD detection
- To overcome these limitations, we proposed the Attention Masking method to improve model to more accurately distinguish between ID and OOD data



Objective of Attention Masking OOD detection

Introduction

- Utilizing Vision Transformer(ViT) Models
- Identifying and preprocessing important patches
 - This approach allows the model to focus more on significant parts of the image
- Applying Mahalanobis distance calculation
- Near Out-of-Distribution image classification
 - Got improved results on CIFAR-10(in) vs. CIFAR-100(out) and vice versa



ENHANCING THE RELIABILITY OF OUT-OF-DISTRIBUTION IMAGE DETECTION IN NEURAL NETWORKS

Shiyu Liang

Coordinated Science Lab, Department of ECE
University of Illinois at Urbana-Champaign
sliang26@illinois.edu

Yixuan Li

University of Wisconsin-Madison*
sharonli@cs.wisc.edu

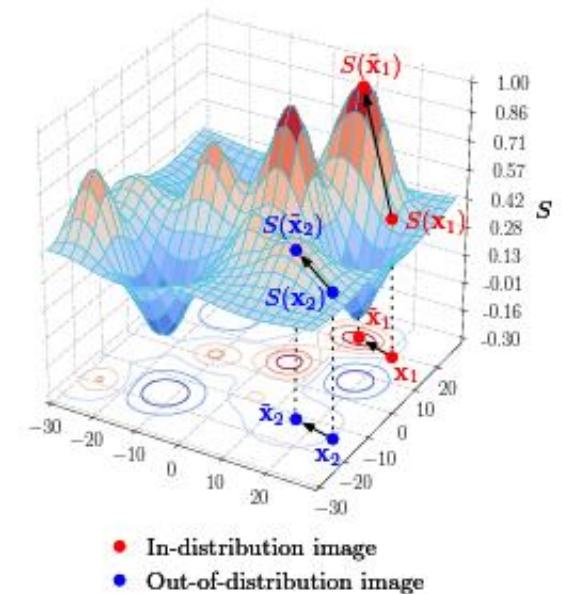
R. Srikanth

Coordinated Science Lab, Department of ECE
University of Illinois at Urbana-Champaign
rsrikanth@illinois.edu

Input pre-processing(ODIN)

Related works

- In contrast to traditional FGSM (Fast Gradient Sign Method), which focuses on reducing the max probability by adding noise
- The approach in the referenced study introduces noise to increase the confidence score.
- The input pre-processing equation used in this method is:
 - $\hat{x} = x - \epsilon \cdot \text{sign}(\nabla_x \log S(x)) = x - \epsilon \cdot \text{sign}(\nabla_x L(x))$
 - We call this method as Reversed FGSM



Attention Masking for Improved Near Out-of-Distribution Image Detection

Minho Sim

School of Computing

KAIST

Daejeon, Republic of Korea

smh3946@kaist.ac.kr

Jongwhoa Lee

School of Computing

KAIST

Daejeon, Republic of Korea

jongwhoa.lee@kaist.ac.kr

Ho-Jin Choi

School of Computing

KAIST

Daejeon, Republic of Korea



hojinc@kaist.ac.kr

CLS Token for Attention Rollout Algorithm

Related works

- **Class Token(CLS Token)**
 - CLS Token is a special token added to the input of the ViT model that learns the summary information of the entire image.
 - In each Transformer layer, the CLS token interacts with other patches through the attention mechanism, ultimately aggregating the important information of the entire image.
- **Attention Weights**
 - Attention rollout heatmaps show how much the CLS token focuses on the image. This is what attention weights represent.

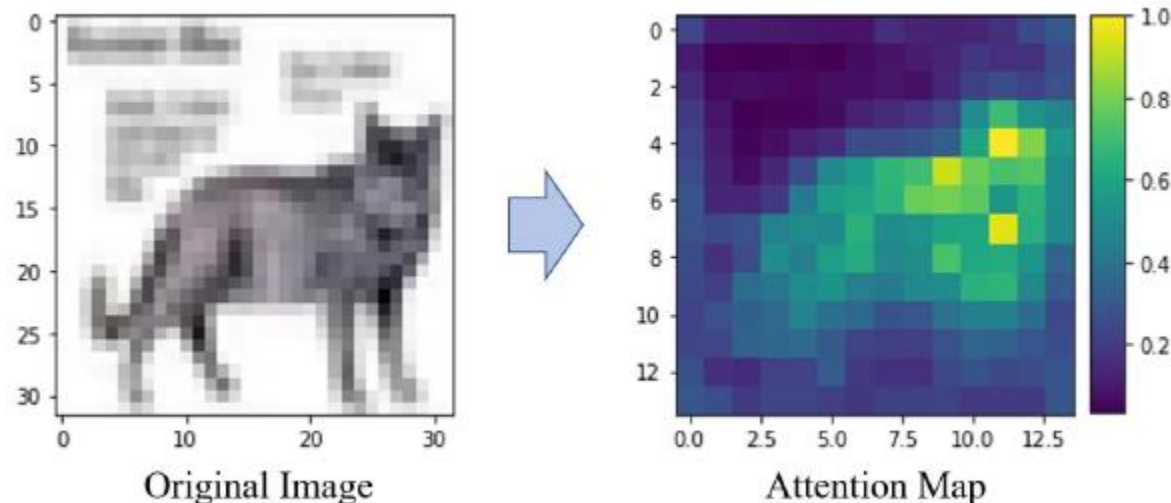
Attention heatmap



Visualization of Attention Mechanism

methods

- Visualizing attention maps helps in understanding the model's decision-making process.
- The attention mechanism allows the model to focus on the most relevant parts of the image.




Masking strategy


Related works

- In this paper the ratio was set to 20%

Masking Strategy	Patches to be Masked
Top-Ratio	Top N% of attention values
Bottom-Ratio	Bottom N% of attention values



(a) top-ratio



(b) bottom-ratio



How can we detect OODs?

- The feature is extracted from the penultimate layer to capture high-level abstract representations of the input
- Calculate the distance between the instance penultimate layer & mean penultimate vector of ID distribution.
 - Close distance: ID
 - Far distance: OOD
- Which distance should we implement?

Mahalanobis distance based OOD detection

Mahalanobis distance is a measure of the distance between a point and a distribution.

- It is used to identify how far a point is from the mean of a distribution, normalized by the covariance of the distribution.
- In the context of OOD detection, it helps in distinguishing between ID and OOD samples based on their distance from the ID distribution.



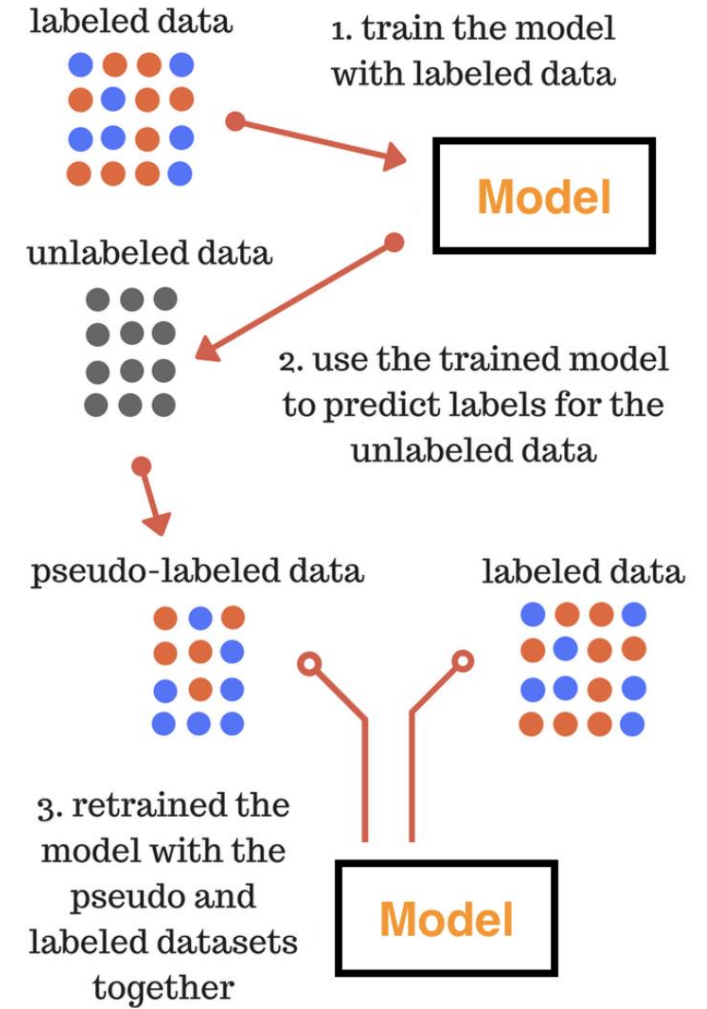
Mahalanobis distance based OOD detection

- $\widehat{\mu}_c = \frac{1}{N_c} \sum_{i: y_i \in c} f(x_i)$, c : class label
- $\widehat{\Sigma} = \frac{1}{N} \sum_c \sum_{i: y_i = c} (f(x_i) - \widehat{\mu}_c)(f(x_i) - \widehat{\mu}_c)^T$
- $M(x) = \min_c (f(x_i) - \widehat{\mu}_c)^T \widehat{\Sigma}^{-1} \cdot (f(x_i) - \widehat{\mu}_c)$
- Attention Masking effectively helps in differentiating ID and OOD samples.
- By focusing on the most relevant parts of the image, the model enhances its ability to distinguish between ID and OOD data
 - 1. Assume a Gaussian distribution for each of the 10 labels(Cifar-10)
 - 2. Calculate the Mahalanobis distance for each distribution
 - 3. Select the nearest label

Pseudo Labeling for input pre-processing

methods

- Pseudo-labeling is applied in our method, where OOD(Out-of-distribution) data is unknown .
- In our method, unknown data is pseudo-labeled with the nearest label.
- $\hat{x} = x - \epsilon \cdot \text{sign}(\nabla_x L(x))$, $L(x)$: CE loss
 - Target is unknown for test dataset



Input pre-processing


methods

- Reduce the distortion of the original image
- So, input pre-processing is applied to partial patches with attention scores strategies.
- The input pre-processing equation used in this method is:
 - $\hat{x} = x - \epsilon \cdot \text{sign}(\nabla_x L(x))$, $L(x)$: CE loss with pseudo label

Pipeline of Attention Adversarial OOD

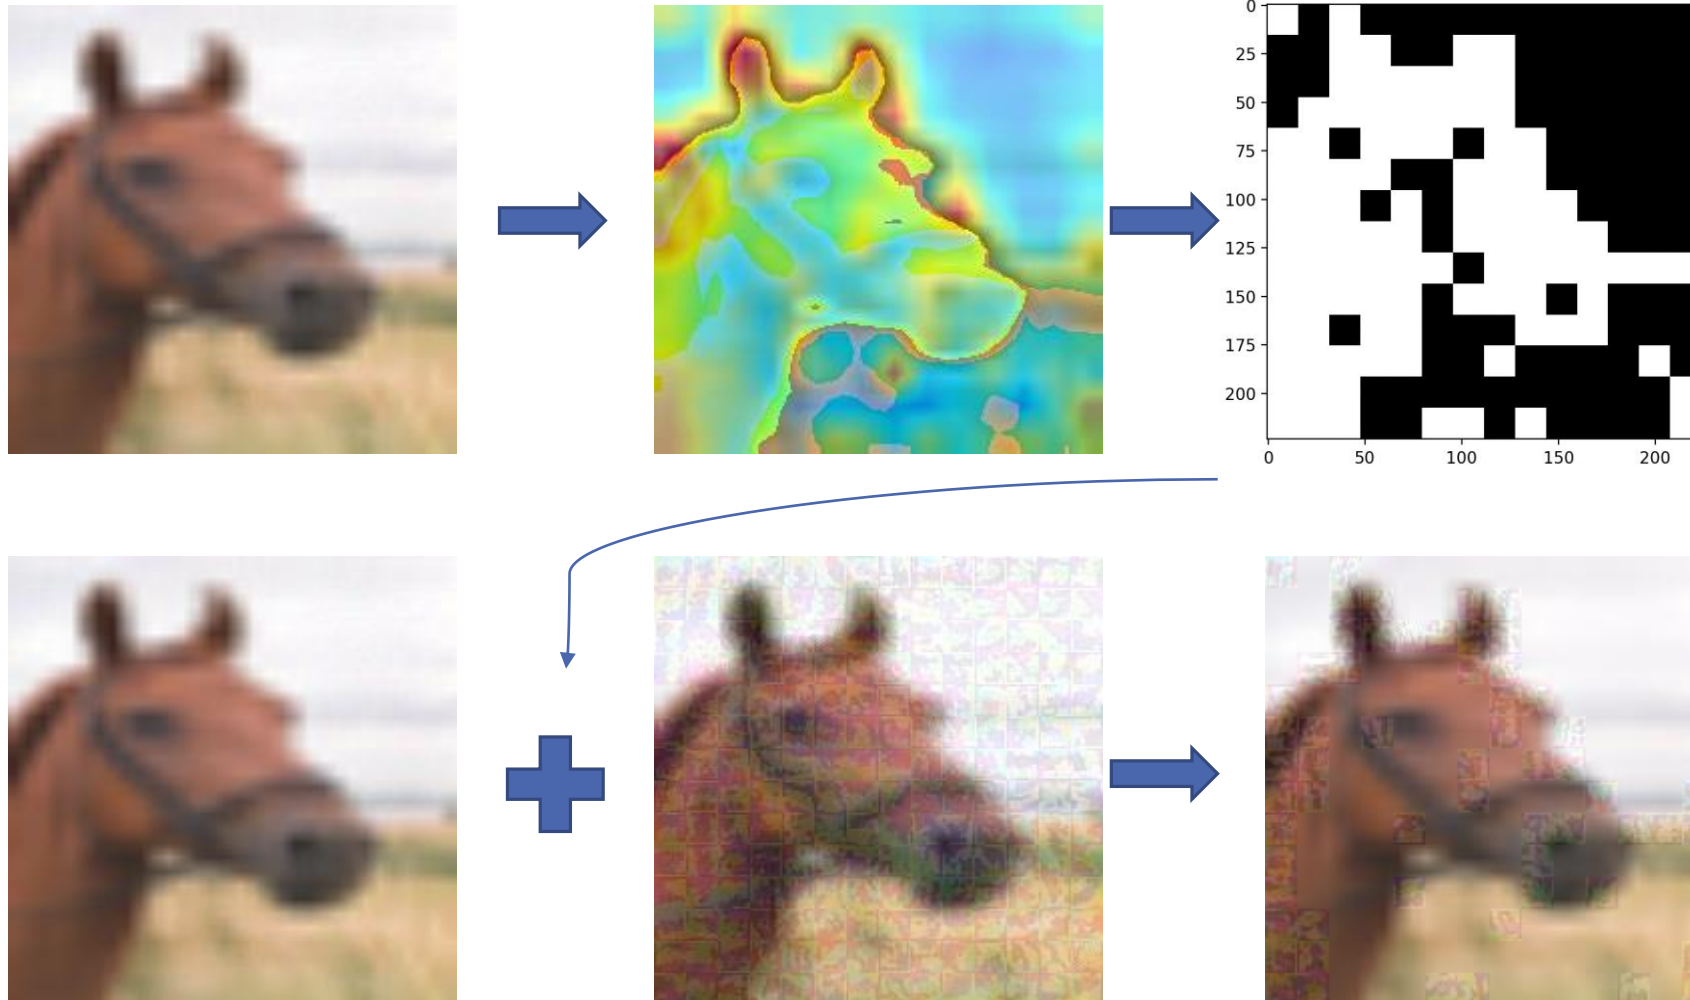
methods

- 1. The process begins with the input image. This is a raw image from the dataset
- 2. The input image is divided into smaller, non-overlapping patches. Each patch is processed individually by the Vision Transformer (ViT) model
- 3. The Vision Transformer model computes attention scores for each patch.
 - An attention map is generated, highlighting the areas of the image that the model focuses on the most.
- 4. Based on the attention scores, more important patches are identified and subtract adversarial noise.
 - This step enhances the model's focus on the significant parts of the image by removing or de-emphasizing the less relevant areas.
- 5. Generate perturbed image
 - Adversarial image by fgsm attack
 - Adversarial image * masking + original image * (1 – masking)



Pipeline of Attention Adversarial OOD

methods



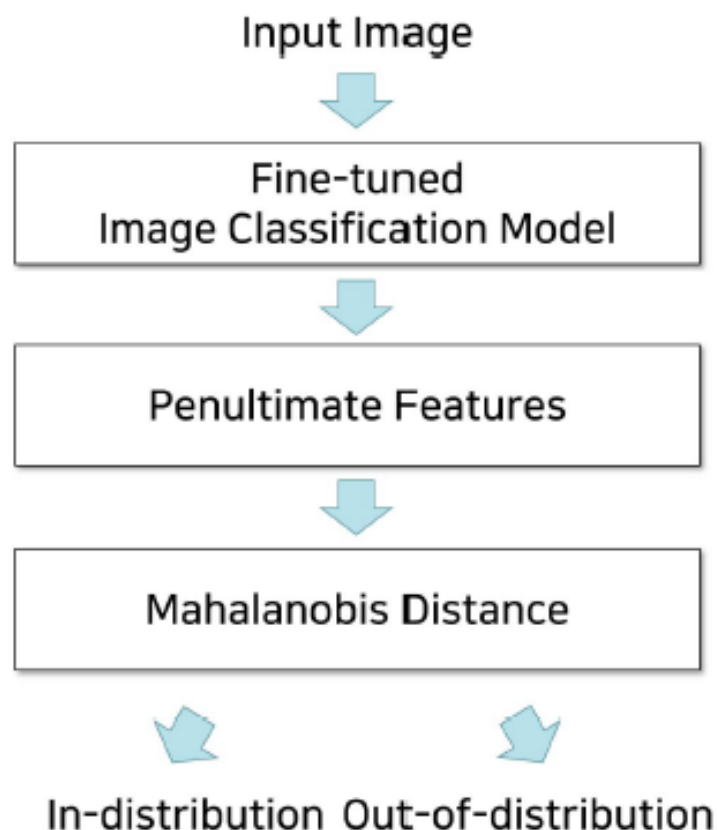
Pipeline of Attention Adversarial OOD

methods

- 6. Input the masked image back into the ViT model to extract new feature vectors(F_{adv}).
- 7. Mahalanobis Distance Calculation and classify ood
 - Mahalanobis distance is calculated using the feature vector extracted from the penultimate layer of the ViT model through which the masked image is passed.
 - Mahalanobis distance is used to measure how much a given feature vector deviates from the training data distribution
 - $D_M = \sqrt{(F_{adv} - \mu)^T \Sigma^{-1} (F_{adv} - \mu)}$

Pseudo code

methods



Algorithm 1 Differential masking OOD detection

Require: Pre-trained ViT model, Input image I , OOD detection score threshold T , Perturbation ϵ

Ensure: OOD Detection score

```
1: Step 1: Extract Attention Weights
2: Initialize the ViT model with pre-trained weights
3: Extract attention weights from all transformer layers for input  $I$ 
4: Step 2: Compute Attention Map
5:  $A_{\text{rollout}_0} = I$  (Identity matrix)
6: for each layer  $L$  in ViT do
7:   Compute attention scores  $A_L$  for each patch in  $I$ 
8:    $A_{\text{rollout}_L} = (A_L + I) \cdot A_{\text{rollout}_{L-1}}$ 
9: end for
10:  $A_{\text{final}} = A_{\text{rollout}_L}$ 
11: Step 3: Identify Important Patches
12: Determine threshold  $T$  for masking based on attention scores
13: for each patch  $p$  in  $I$  do
14:   if Attention score of  $p > T$  then
15:     Point patch  $p$ 
16:   end if
17: end for
18: Step 4: Adversarial Noise
19: Set parameters:  $\epsilon$ , iterations,  $\tau$ 
20: for each image  $I$ , label in dataset do
21:   Initialize image gradient
22:   for 1 to iterations do
23:     Compute loss and gradients
24:     Apply attention filter:  $\text{grad} \leftarrow \text{grad} \times (\text{attention} > \tau)$ 
25:     Update image:  $I \leftarrow I - \frac{\epsilon}{\text{iterations}} \cdot \text{sign}(\text{grad})$  for pointed patches only
26:     Clamp image to valid range
27:   end for
28:   Evaluate model accuracy with and without partial input preprocessing
29: end for
30: Step 5: Evaluate Updated Image
31: Forward perturbed image  $I_{\text{perturbed}}$  through ViT model
32: Extract penultimate layer features  $F_{\text{perturbed}}$ 
33: Calculate Mahalanobis distance  $D_M$  using  $F_{\text{perturbed}}$  and training distribution statistics
34: Step 6: OOD Detection
35: Determine OOD score based on  $D_M$ 
36: if  $D_M > \text{threshold}$  then
37:   Classify  $I$  as OOD
38: else
39:   Classify  $I$  as ID
40: end if
41: return OOD Detection score
```

Experiment

experiment

- **Datasets:** CIFAR-10(id), CIFAR-100(ood)
- **Models:** CIFAR-10-finetuned ViT
- **Image Processing:**
 - CIFAR images scaling 32x32 to 224x224
- **Metrics:** AUROC, AUPR_ood
- **Experiment:**
 - Near-OOD detection with fine-tuning
- **Baseline method:**
 - Mahalanobis Distance

Result

experiment

Table 1: Perturb Attention Top 50%



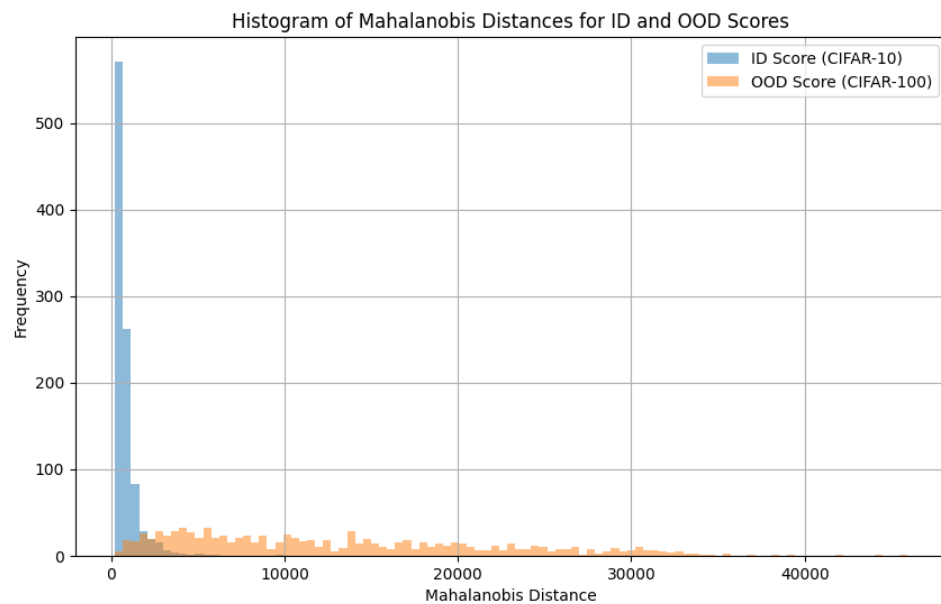
Metric	baseline	$\epsilon = 0.1$	$\epsilon = 0.01$	$\epsilon = 0.001$	$\epsilon = -0.001$	$\epsilon = -0.01$
auroc	0.9857	0.9371	0.9379	0.9846	0.9859	0.9348
aupr	0.9875	0.9465	0.945	0.9865	0.9877	0.9463

Table 2: Perturb Attention Bottom 50%

Metric	baseline	$\epsilon = 0.1$	$\epsilon = 0.01$	$\epsilon = 0.001$	$\epsilon = -0.001$	$\epsilon = -0.01$
auroc	0.9857	0.8607	0.8907	0.9844	0.9860	0.8544
aupr	0.9875	0.8782	0.9008	0.9846	0.9880	0.8991

baseline

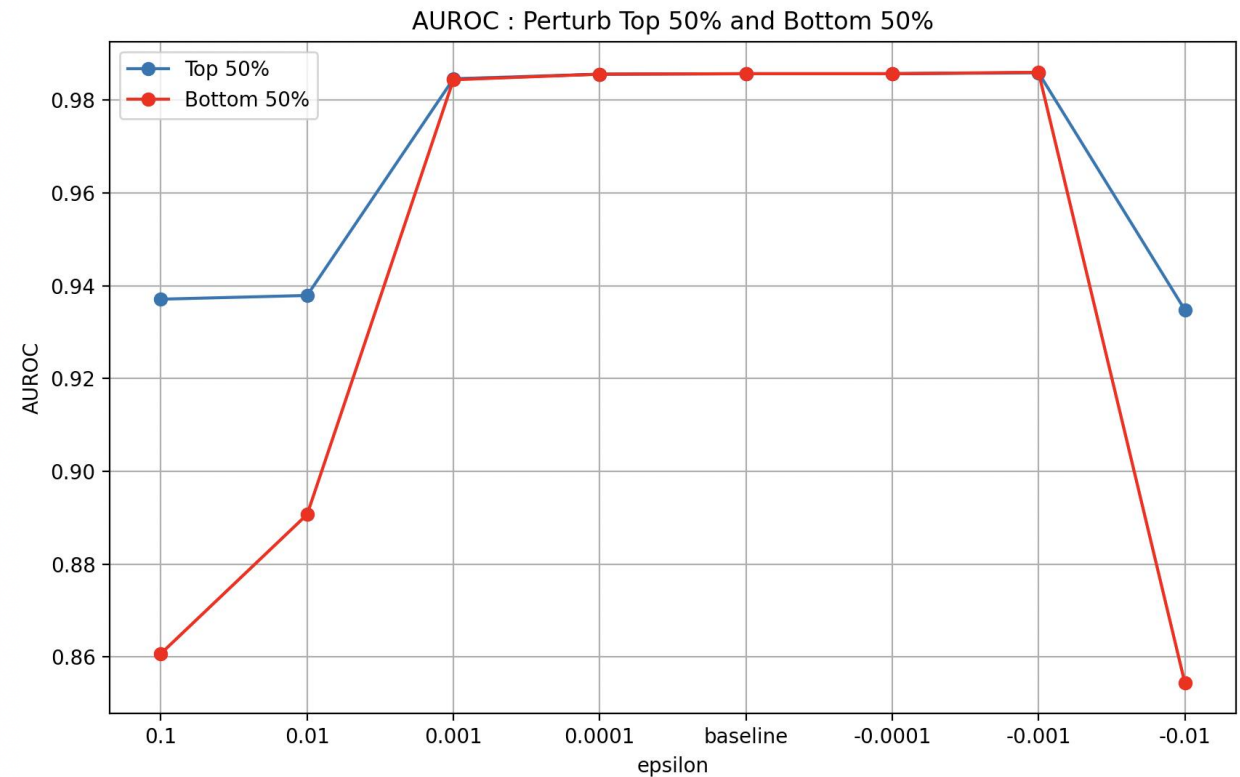
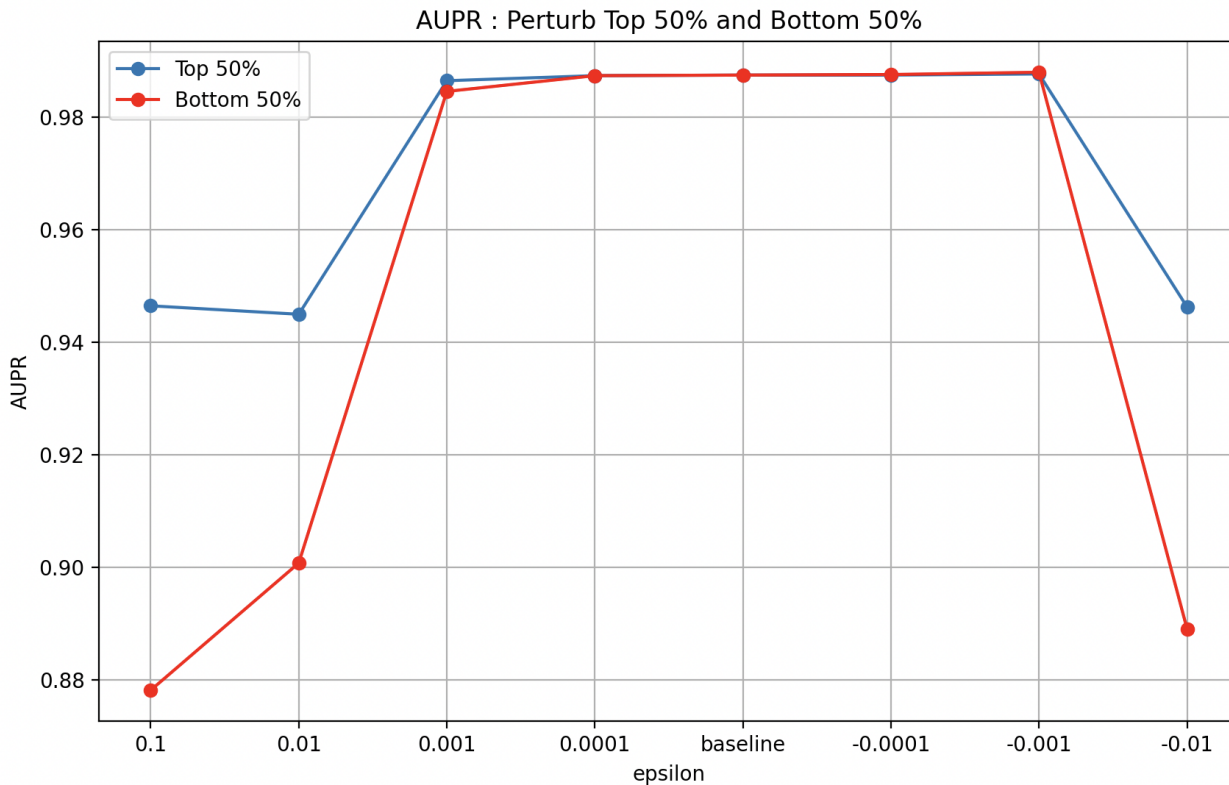
experiment



Result


experiment

- Compare two strategy top 50% and bottom 50%




Result

Baseline



$\epsilon = 0.1$



Limitation


- **Limitation of method**
 - Dataset is too easy, making it difficult to confirm whether detection is being performed accurately. So we should try using a different dataset.
 - The base performance of ViT is too high
- Result may depend on fine-tuned ViT model

Appendix

CLS Token for Attention Rollout Algorithm

appendix

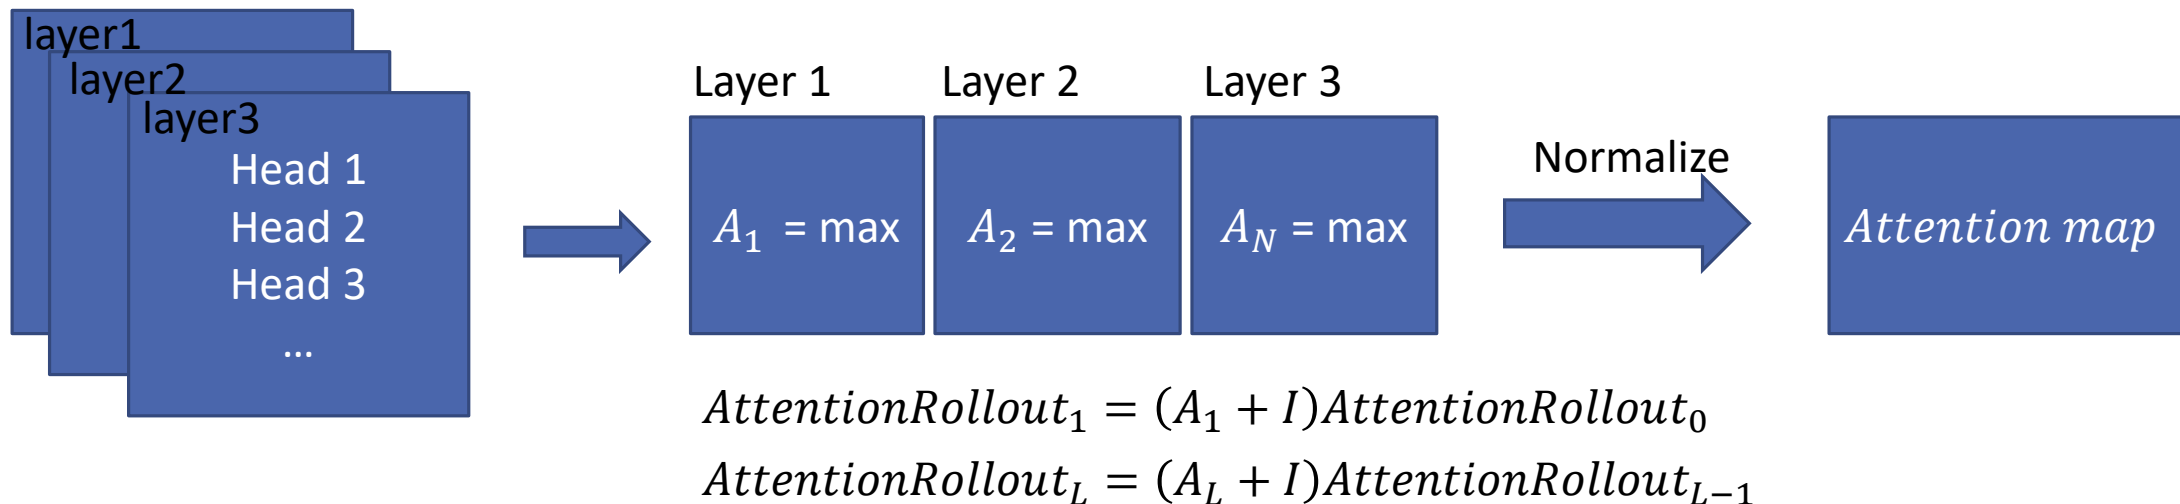
- **Calculation of Attention Rollout**
 - Attention Rollout calculates the degree to which the CLS token attends to the patches at each layer.
 - To do this, it sequentially multiplies the attention weights of each layer to compute the final attention map.
- **Conclusion**
 - The Attention Rollout Algorithm calculates how much the CLS token attends to each patch to generate the final attention map.
 - This final map visually expresses the importance of each patch in the model's final output.
 - In other words, it calculates and visualizes the importance of patches based on the attention given by the CLS token to each patch.



Attention rollout algorithm

appendix

- Proposed at “Quantifying Attention Flow in Transformers”
 - Visualizes and interprets how attention mechanisms in transformers contribute to model decisions.
 - Visualizes the effect of multi-head attention in a single graph.
- Attention rollout Visualize multi-head Attention`s effect by one graph



Attention rollout algorithm

appendix

- **Process of Attention rollout**
 - **1. Extract Attention Weights:** $A_L = \text{Attention Weights from Layer } L$
 - For each transformer layer, extract the attention weights for all attention heads.
 - **2. Head Fusion:** $A_L = \frac{1}{h} \sum_{k=1}^h A_L^k$ or $A_L = \max_{k=1}^h A_L^k$
 - Integrate the attention weights from multiple heads into a single matrix by averaging or taking the maximum values.
 - **3. Recursive Aggregation:** $\text{AttentionRollout}_L = (A_L + I) \text{AttentionRollout}_{L-1}$
 - Here, A_{ij} represents how much the j -th patch in the previous layer attends to the i -th patch in the current layer.
 - $\text{AttentionRollout}_0$ is identity matrix
 - Calculate the aggregated attention rollout using the attention weights.
 - Starting from the final layer, recursively multiply the attention weights through the layers to aggregate the overall attention for each token.
 - **4. Normalize Attention Scores:** $\text{NormalizedAttention}_L = \frac{\text{AttentionRollout}_L}{\sum_{k=1}^N \text{AttentionRollout}_{ij}}$
 - Normalize the aggregated attention scores to ensure they sum to 1, providing a clear distribution for focus
 - **5. Visualization**
 - Visualize the result attention map